


Probabilistic Watershed: Sampling all spanning forests for seeded segmentation and semi-supervised learning

Enrique Fita Sanmartín, Sebastian Damrich, Fred A. Hamprecht
Supervisor: Prof. Fred A. Hamprecht


HCI/IWR at Heidelberg University

Notation: edge cost vs edge weight

$$w(e) = \exp(-\mu c(e)), \mu \geq 0$$



(a) Graph's edge-costs




(b) Graph's edge-weights, $\mu = 1$

$$c(G) = \sum_{e \in E_G} c(e)$$


$$w(G) = \prod_{e \in E_G} w(e)$$

Watershed computes a minimum cost spanning forest (mSF)


The assignment of q is doubtful



(a) Graph with two seeds



(b) mSF: Cost=1.32



(c) SF: Weight=1.4

Figure: (2a) Graph with edge-costs and two seeds. (2b) The mSF (Watershed segmentation) assigns q to seed s_1 . (2c) Low cost, but not minimum, spanning forest connecting q to s_2 . The Watershed has a winner-takes-all behaviour. It only considers the mSF while ignores the rest of seed separating spanning forests

What is the probability of sampling a forest such that a node of interest is assigned to a certain seed?

Gibbs Distribution of \mathcal{F}^{s_2}

$$\mathcal{F}_{s_1}^{s_2} := \{2\text{-forests separating } s_1 \text{ and } s_2\}$$

Definition: probability distribution over the forests

Given $J \in \mathbb{R}_{\geq 0}$

$$\Pr^* = \arg \min_{\Pr} \sum_{f \in \mathcal{F}_{s_1}^{s_2}} \Pr(f) c(f), \quad \text{s.t.} \quad \sum_{f \in \mathcal{F}_{s_1}^{s_2}} \Pr(f) = 1 \quad \text{and} \quad \mathcal{H}(\Pr) = J, \quad (1)$$

where $\mathcal{H}(\Pr)$ = entropy \Pr

Solution: Gibbs Probability distribution

$$\Pr^*(f) = \frac{\exp(-\mu c(f))}{\sum_{f' \in \mathcal{F}_{s_1}^{s_2}} \exp(-\mu c(f'))} = \frac{w(f)}{\sum_{f' \in \mathcal{F}_{s_1}^{s_2}} w(f')}$$

$\mathcal{F}_{s_1}^{s_2} := \{2\text{-forests separating } s_1 \text{ and } s_2\}$

$\mathcal{F}_{s_1, q}^{s_2} := \{f \in \mathcal{F}_{s_1}^{s_2} : s_1 \text{ and } q \text{ are connected}\}$

Probability q is connected with s_1

$$\Pr(q \sim s_1) := \frac{w(\mathcal{F}_{s_1, q}^{s_2})}{w(\mathcal{F}_{s_1}^{s_2})} = \frac{\sum_{f \in \mathcal{F}_{s_1, q}^{s_2}} w(f)}{\sum_{f \in \mathcal{F}_{s_1}^{s_2}} w(f)}.$$

Example

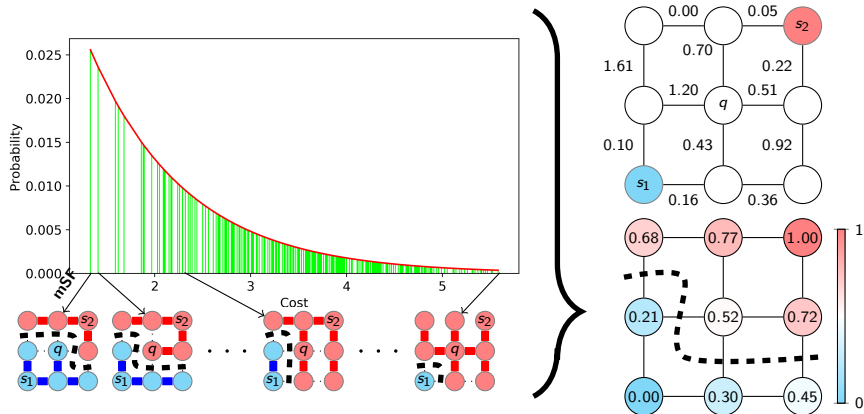


Figure: (Top left) Gibbs distribution over all spanning forests. (Bottom right) Probabilistic Watershed probabilities for assigning a node to s_2 . The Probabilistic Watershed computes the expected seed assignment of every node for a Gibbs distribution over all exponentially many spanning forests in closed-form. It thus avoids the winner-takes-all behaviour of the Watershed.

$\mathcal{F}_{s_1}^{s_2} := \{2\text{-forests separating } s_1 \text{ and } s_2\}$

$\mathcal{F}_{s_1,q}^{s_2} := \{f \in \mathcal{F}_{s_1}^{s_2} : s_1 \text{ and } q \text{ are connected}\}$

Probability q is connected to s_1

$$\Pr(q \sim s_1) := \frac{w(\mathcal{F}_{s_1,q}^{s_2})}{w(\mathcal{F}_{s_1}^{s_2})} = \frac{\sum_{f \in \mathcal{F}_{s_1,q}^{s_2}} w(f)}{\sum_{f \in \mathcal{F}_{s_1}^{s_2}} w(f)}.$$

Question: How do we compute $\mathcal{F}_{s_1}^{s_2}$ and $\mathcal{F}_{s_1,q}^{s_2}$?

Answer: Matrix Tree Theorem (MTT).

Matrix Tree Theorem

Matrix Tree Theorem (MTT)


For any weighted multigraph G , the sum of the weights of the spanning trees of G , $w(\mathcal{T})$, is equal to

$$w(\mathcal{T}) := \sum_{t \in \mathcal{T}} \prod_{e \in E_t} w(e) = \frac{1}{|V|} \det(L + \frac{1}{|V|} \mathbb{1} \mathbb{1}^\top) = \det(L^{[v]}),$$


where $\mathbb{1}$ is the column full of 1. $L^{[v]}$ is the matrix obtained after removing the row and column corresponding to an arbitrary node v .

Computation of $w(\mathcal{F}_{s_1}^{s_2})$

The relation between forests and trees allows the use of the MTT.



(a) Forest.



(b) Tree


Computation of $w(\mathcal{F}_{s_1}^{s_2})$

r_{uv}^{eff} = effective resistance distance between the nodes u and v .


Theorem (Consequence of the MTT)

$$w(\mathcal{F}_{s_1}^{s_2}) = w(\mathcal{T}) r_{s_1 s_2}^{\text{eff}} \propto r_{s_1 s_2}^{\text{eff}}.$$

Computation of $w(\mathcal{F}_{s_1, s_2}^q)$ and $w(\mathcal{F}_{s_2, s_1}^q)$




(a) spanning tree $t \in \mathcal{T}$




(b) forest $f \in \mathcal{F}_{s_1}^{s_2}$

Figure: $\mathcal{F}_{s_1, q}^{s_2}$ and $\mathcal{F}_{s_2, q}^{s_1}$ form a partition of $\mathcal{F}_{s_1}^{s_2}$ since q must be connected either to s_1 or s_2 , but not to both. Given an spanning tree (dashed lines represent all possible spanning trees), to form a forest in $\mathcal{F}_{s_1}^{s_2}$ we will have to remove an edge from the red part, forming a forest in $\mathcal{F}_{s_2, q}^{s_1}$, or from the blue part, forming a forest in $\mathcal{F}_{s_1, q}^{s_2}$.


Computation of $w(\mathcal{F}_{s_1, s_2}^q)$ and $w(\mathcal{F}_{s_2, s_1}^q)$




(a) spanning tree $t \in \mathcal{T}$



(b) forest $f \in \mathcal{F}_{s_1}^{s_2}$



(c) forest $f \in \mathcal{F}_{s_2}^q$



(d) forest $f \in \mathcal{F}_{s_1}^q$

Figure: Analogously \mathcal{F}_{s_1, s_2}^q and $\mathcal{F}_{q, s_2}^{s_1}$ form a partition of $\mathcal{F}_{s_1}^q$, and \mathcal{F}_{s_2, s_1}^q and $\mathcal{F}_{q, s_1}^{s_2}$ form a partition of $\mathcal{F}_{s_2}^q$ respectively.

Computation $w(\mathcal{F}_{s_1, s_2}^q), w(\mathcal{F}_{s_2, s_1}^q)$


Linear system with 3 unknowns and 3 equations.

$$\begin{aligned}w(\mathcal{F}_{s_1, q}^{s_2}) + w(\mathcal{F}_{s_1, s_2}^q) &= w(\mathcal{F}_{s_2}^q) \\w(\mathcal{F}_{s_1, s_2}^q) + w(\mathcal{F}_{s_2, q}^{s_1}) &= w(\mathcal{F}_{s_1}^q) \\w(\mathcal{F}_{s_1, q}^{s_2}) + w(\mathcal{F}_{s_2, q}^{s_1}) &= w(\mathcal{F}_{s_1}^{s_2}).\end{aligned}$$

Probabilistic Watershed corresponds to the effective resistance distance

Probability measures the gap of the triangle inequality

$$\Pr(q \sim s_1) = \frac{w(\mathcal{F}_{s_2}^q) + w(\mathcal{F}_{s_1}^{s_2}) - w(\mathcal{F}_{s_1}^q)}{2w(\mathcal{F}_{s_1}^{s_2})} = \frac{r_{s_2 q}^{\text{eff}} + r_{s_2 s_1}^{\text{eff}} - r_{s_1 q}^{\text{eff}}}{2r_{s_2 s_1}^{\text{eff}}}. \quad (2)$$



$$\Pr(q \sim s_1) \geq \Pr(q \sim s_2) \iff r_{s_1 q}^{\text{eff}} \leq r_{s_2 q}^{\text{eff}}.$$

Theorem

The probability, $x_q^{s_2}$, that a Random Walker starting at node q reaches first s_2 before reaching s_1 ([Grady, 2006]) is equal to the probability defined by the Probabilistic Watershed.

$$x_q^{s_2} = P(q \sim s_2).$$

Power Watershed minimization problem

$$x^* = \arg \min_x \sum_{e=(v,u) \in E} (w(e))^\alpha (|x_v - x_u|)^\beta, \text{ s.t. } x_{s_1} = 1, x_{s_2} = 0, \quad (3)$$

Random Walker: $\alpha = 1, \beta = 2$

Power Watershed: $\alpha \rightarrow \infty, \beta = 2$

α and μ are equivalent.

$$w(e) = \exp(-\mu c(e)) \rightarrow (w(e))^\alpha = \exp(-\mu \alpha c(e)) = \exp(-\mu_\alpha c(e))$$

Power Watershed counts minimum cost/maximum weight spanning forests

Given $S = \{s_1, s_2\}$ a set of seeds, let us denote the potential of node q being assigned to seed s_1 by the Power Watershed with $\beta = 2$ as $x_q^{\text{PW}_1}$. Let further w_{\max} be $\max_{f \in \mathcal{F}_{s_1}^{s_2}} w(f)$. Then

$$x_q^{\text{PW}_1} = \Pr_{\infty}(q \sim s_1) := \frac{|\{f \in \mathcal{F}_{s_1, q}^{s_2} : w(f) = w_{\max}\}|}{|\{f \in \mathcal{F}_{s_1}^{s_2} : w(f) = w_{\max}\}|}.$$

Power Watershed counts mSF




Figure: Forest-interpretation of Power Watershed. The Power Watershed computes the ratio between the mSFs connecting a node to s_2 and all possible mSFs. (7f) indicates the allowed Random Walker transitions when $\mu \rightarrow \infty$ with headed arrows. The Random Walker interpretation of the Power Watershed breaks down in the limit case since a Random Walker starting at node q does not reach any seed, but oscillates along the bold arrow.

- Probabilistic Watershed = Random Walker
- Random Walker/Probabilistic Watershed probabilities are proportional to the triangle inequality gap.

$$\Pr(q \sim s_1) \propto r_{s_2 q}^{\text{eff}} + r_{s_2 s_1}^{\text{eff}} - r_{s_1 q}^{\text{eff}}$$

- Power Watershed counts minimum cost spanning forests



Couprie, C., Grady, L., Najman, L., and Talbot, H. (2011).
Power watershed: A unifying graph-based optimization framework.
IEEE Transactions on Pattern Analysis and Machine Intelligence,
33:1384–1399.



Grady, L. (2006).
Random walks for image segmentation.
IEEE Transactions on Pattern Analysis and Machine Intelligence.

The End